

House Price Analysis

Authors: Ricky Wong

Overview

Jim's Real Estate wants to help homeowners make the most of their property and would like to know how renovations affect house prices. Using linear regression we found two variables that affected house prices the most. Using that information, real estate agents can provide meaningful advice for homeowners.

Business Understanding

There are a variety of variables when determining the value of a property. Creating a model for real estate agents will help them advise homeowners about how renovating affects the value of their property. To do this we will explore house data from King County to build a basic model then reiterate the process to improve on the model.

```
In [1]: #import packages
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import statsmodels.api as sm
from sklearn import preprocessing
import scipy.stats as stats
import seaborn as sns
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

pd.options.mode.chained_assignment = None

%matplotlib inline
```

```
In [2]: #read data from file
kc_house_data df = pd.read_csv('data/kc_house_data.csv')
```

Data Understanding

We have obtained house data from King County to help us build the model. This includes useful information like the price of the property, the living area in square feet and the year the property was renovated to name a few.

After getting an idea of the data we are dealing with we can see most of our data are numerical with some columns missing data like 'waterfront' and 'yr_renovated'. We will need to clean those up later

There are some data which is not really relevant to our model which we can remove.

In [3]: `kc_house_data_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21597 non-null  int64
1   date                  21597 non-null  object
2   price                 21597 non-null  float64
3   bedrooms              21597 non-null  int64
4   bathrooms             21597 non-null  float64
5   sqft_living           21597 non-null  int64
6   sqft_lot              21597 non-null  int64
7   floors                21597 non-null  float64
8   waterfront            19221 non-null  float64
9   view                  21534 non-null  float64
10  condition              21597 non-null  int64
11  grade                  21597 non-null  int64
12  sqft_above             21597 non-null  int64
13  sqft_basement          21597 non-null  object
14  yr_built                21597 non-null  int64
15  yr_renovated            17755 non-null  float64
16  zipcode                21597 non-null  int64
17  lat                    21597 non-null  float64
18  long                   21597 non-null  float64
19  sqft_living15           21597 non-null  int64
20  sqft_lot15             21597 non-null  int64
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB
```

In [4]: `drop = ['date', 'view', 'zipcode', 'lat', 'long', 'sqft_living15', 'sqft_lot15',
kc_house_data_df.drop(drop, axis=1, inplace=True)`

Taking a closer look we can observe a few points

- A possible outlier with 33 bedrooms when the mean is only 3.
- Condition of the house is form 1-5
- There is a grading system from 3 - 13
- And house data from 1900 - 2015
- sqft_basement has '?' which could explain why it is an object instead of being numeric

In [5]: `kc_house_data_df.describe()`

Out[5]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot
count	2.159700e+04	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04
mean	4.580474e+09	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04
std	2.876736e+09	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04
min	1.000102e+06	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02
25%	2.123049e+09	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04

```
In [6]: kc_house_data_df.head(10)
```

Out[6]:

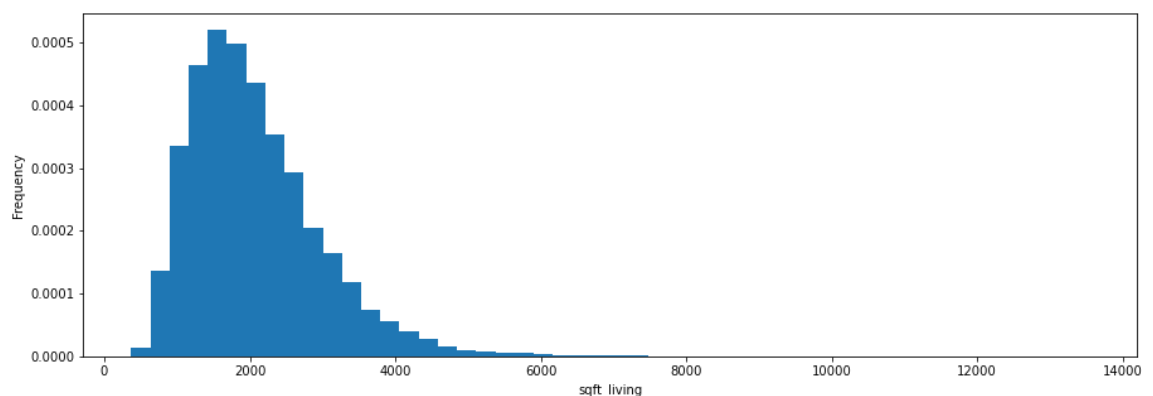
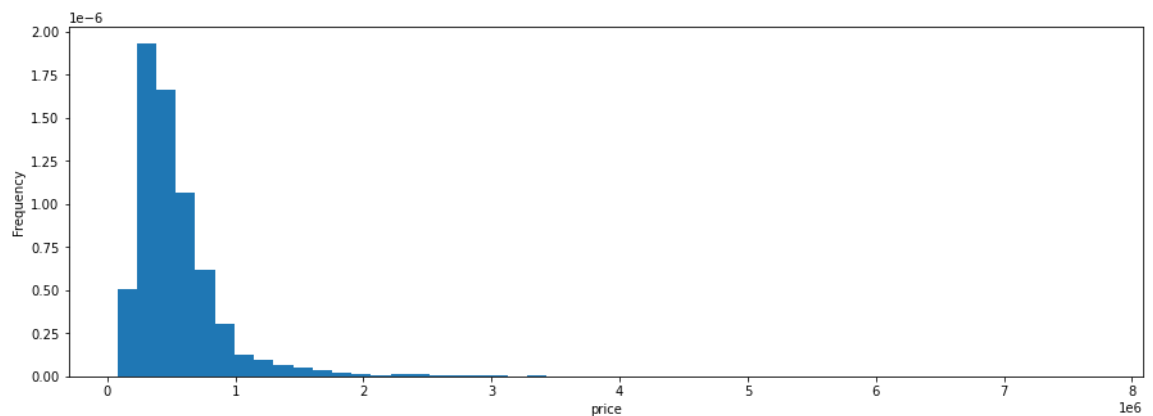
	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	condition	grad
0	7129300520	221900.0	3	1.00	1180	5650	1.0	3	
1	6414100192	538000.0	3	2.25	2570	7242	2.0	3	
2	5631500400	180000.0	2	1.00	770	10000	1.0	3	
3	2487200875	604000.0	4	3.00	1960	5000	1.0	5	
4	1954400510	510000.0	3	2.00	1680	8080	1.0	3	
5	7237550310	1230000.0	4	4.50	5420	101930	1.0	3	1
6	1321400060	257500.0	3	2.25	1715	6819	2.0	3	
7	2008000270	291850.0	3	1.50	1060	9711	1.0	3	
8	2414600126	229500.0	3	1.00	1780	7470	1.0	3	
9	3793500160	323000.0	3	2.50	1890	6560	2.0	3	

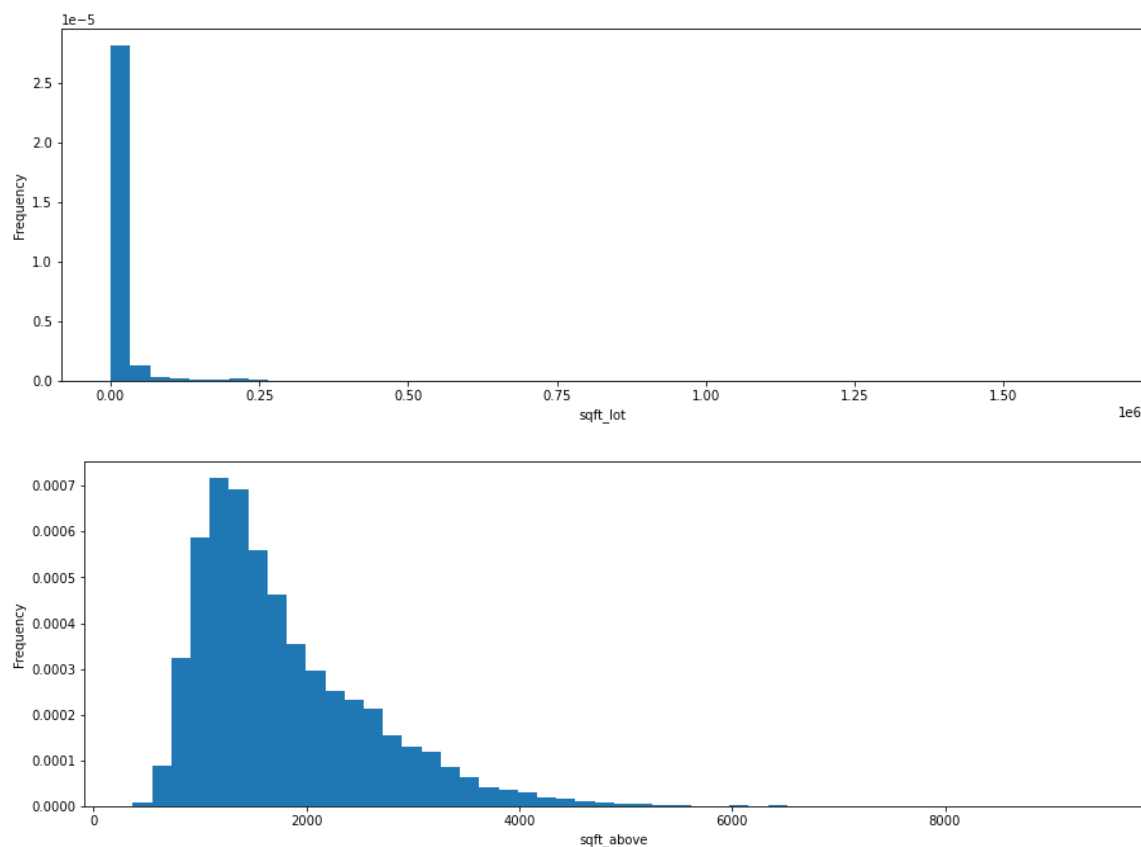
Continuous data

The histograms show the data being positively skewed. We will need to run log transformations to make them more normal.

```
In [7]: continuous = ['price', 'sqft_living', 'sqft_lot', 'sqft_above']

for cont in continuous:
    plt.figure(figsize=(15,5))
    kc_house_data_df[cont].plot.hist(bins=50, density=True)
    plt.xlabel(cont)
    plt.show()
```





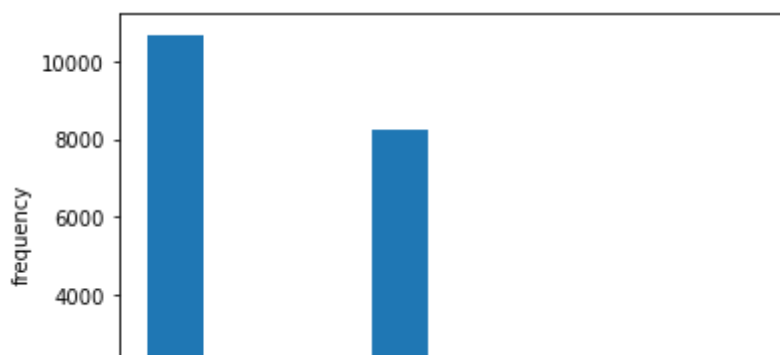
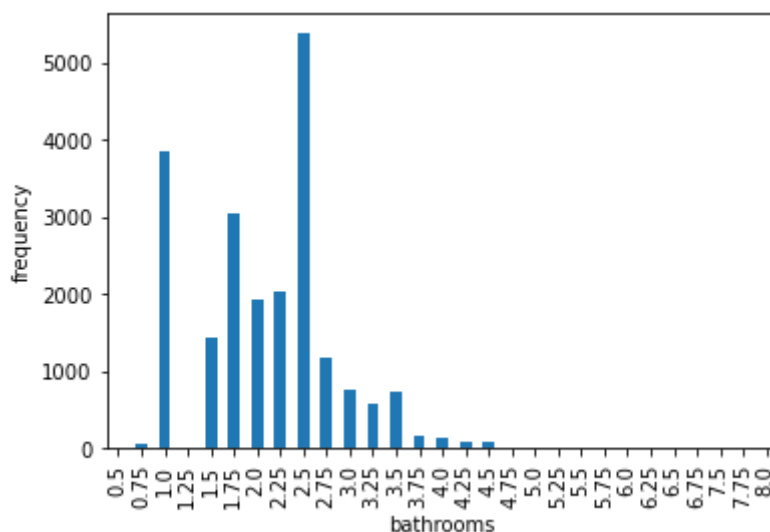
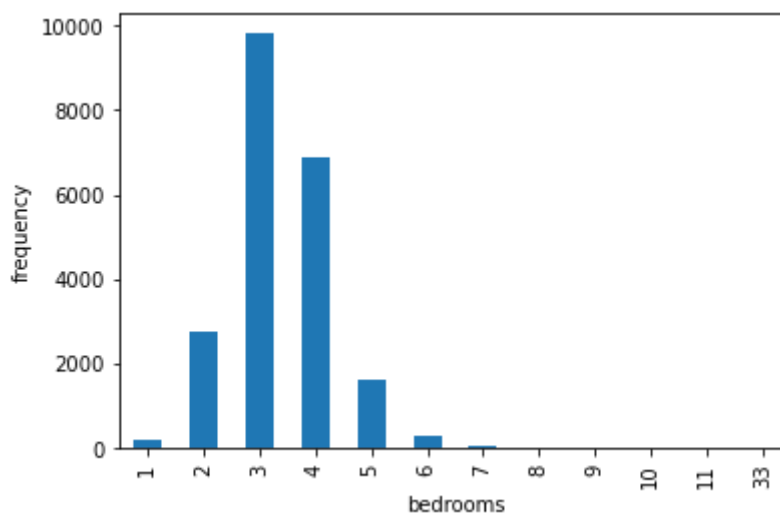
Categorical Data

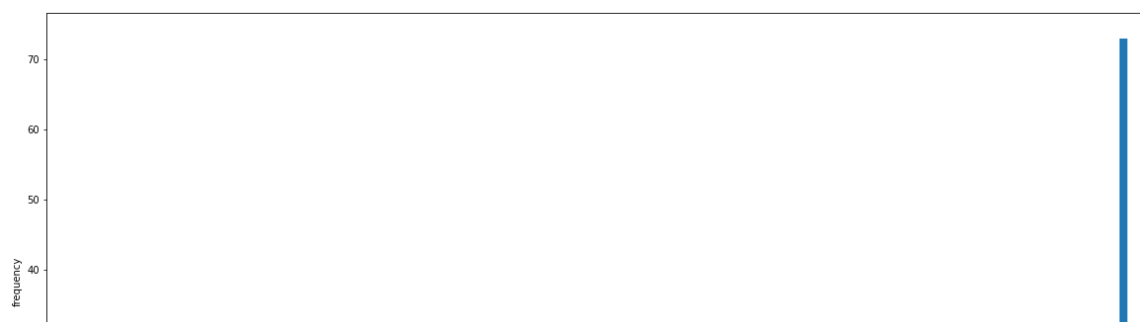
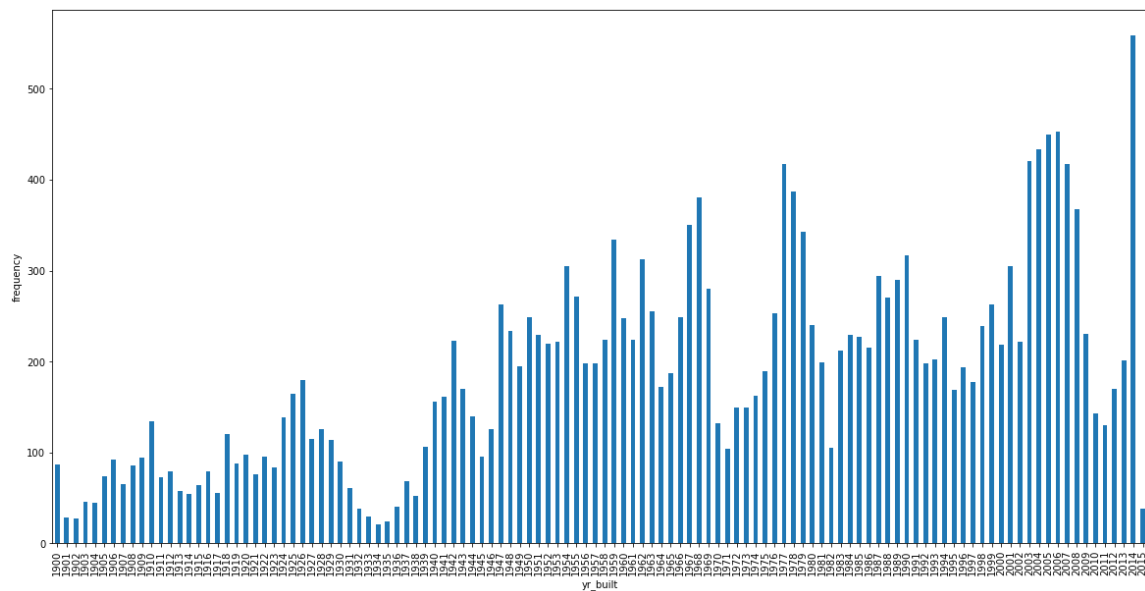
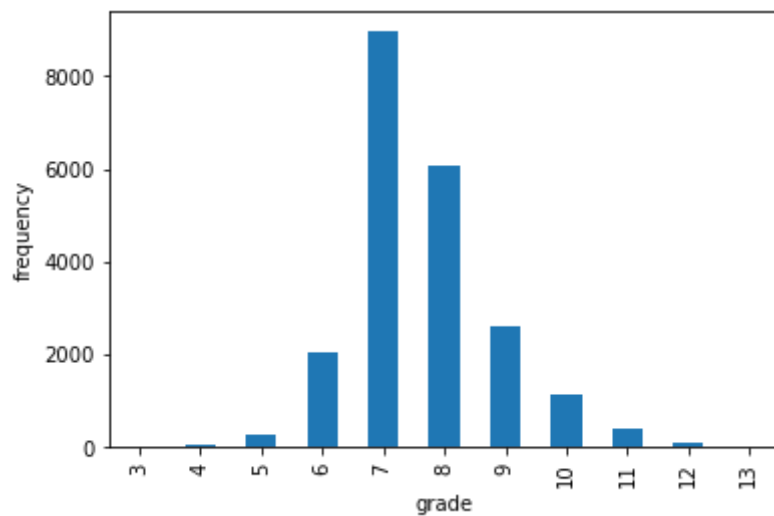
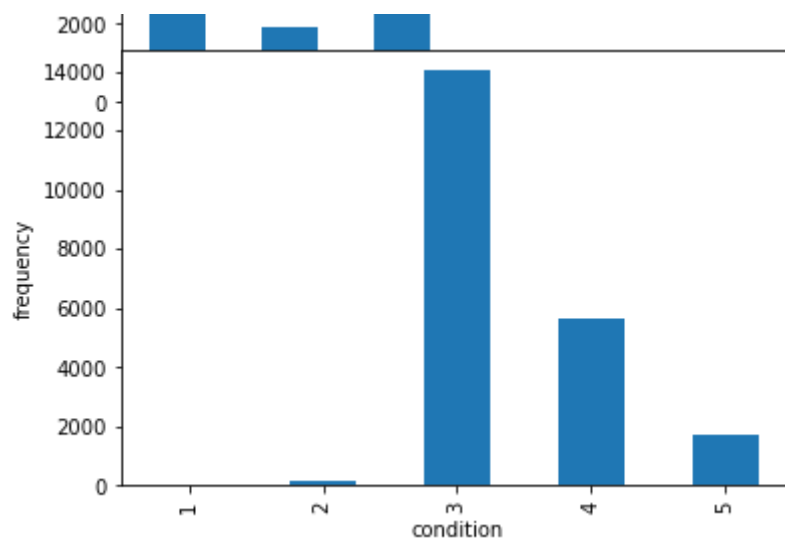
The house with 33 bedrooms is an outlier which we can drop when cleaning
We can see there wasn't many renovations happening (less than 10) until after 1982
Knowing if the property has been renovated or not would be more useful than the year the property was renovated

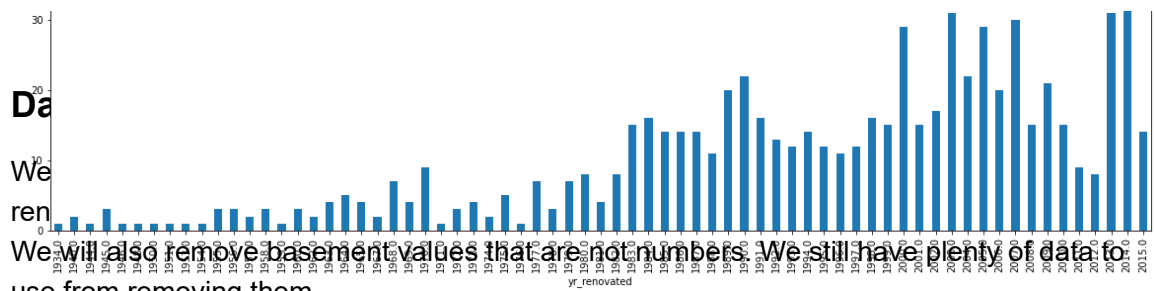
```
In [8]: categorical = ['bedrooms', 'bathrooms', 'floors', 'condition', 'grade', 'yr_built']

for cat in categorical:
    if cat == 'yr_built':
        plt.figure(figsize=(20,10))
        kc_house_data_df.groupby([cat])[cat].count().plot.bar()
        plt.ylabel('frequency')
        plt.show()

    #better view of when properties were renovated
    plt.figure(figsize=(20,10))
    yr_reno = kc_house_data_df.loc[kc_house_data_df['yr_renovated'] > 0]
    yr_reno.groupby(yr_reno['yr_renovated'])['yr_renovated'].count().plot.bar()
    plt.ylabel('frequency')
    plt.show()
```







We will also remove basement values that are not numbers. We still have plenty of data to use from removing them.

Remove the outlier in bedrooms

```
In [9]: clean_house_data_df = kc_house_data_df.dropna(subset=['yr_renovated'])
clean_house_data_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 17755 entries, 0 to 21596
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              17755 non-null  int64
1   price           17755 non-null  float64
2   bedrooms        17755 non-null  int64
3   bathrooms       17755 non-null  float64
4   sqft_living     17755 non-null  int64
5   sqft_lot        17755 non-null  int64
6   floors          17755 non-null  float64
7   condition       17755 non-null  int64
8   grade           17755 non-null  int64
9   sqft_above      17755 non-null  int64
10  sqft_basement   17755 non-null  object
11  yr_built        17755 non-null  int64
12  yr_renovated    17755 non-null  float64
dtypes: float64(4), int64(8), object(1)
memory usage: 1.9+ MB
```

```
In [10]: # Convert basement to numeric and remove the ones that are unknown
clean_house_data_df['sqft_basement'] = pd.to_numeric(clean_house_data_df['sqft_basement'], errors='coerce')
clean_house_data_df.dropna(subset=['sqft_basement'], inplace=True)
clean_house_data_df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 17389 entries, 0 to 21596
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    17389 non-null  int64
1   price                 17389 non-null  float64
2   bedrooms              17389 non-null  int64
3   bathrooms             17389 non-null  float64
4   sqft_living           17389 non-null  int64
5   sqft_lot              17389 non-null  int64
6   floors                17389 non-null  float64
7   condition             17389 non-null  int64
8   grade                 17389 non-null  int64
9   sqft_above            17389 non-null  int64
10  sqft_basement         17389 non-null  float64
11  yr_built              17389 non-null  int64
12  yr_renovated          17389 non-null  float64
dtypes: float64(5), int64(8)
memory usage: 1.9 MB
```

```
In [11]: # remove outlier in bedrooms
clean_house_data_df.loc[(clean_house_data_df['bedrooms'] >= 10)].sort_values('bedrooms')
clean_house_data_df.drop(labels=[15856], axis=0, inplace=True)
```

```
In [12]: #change year to property been renovated or not
clean_house_data_df.loc[(clean_house_data_df['yr_renovated'] > 0), 'yr_renovated'] = 1
clean_house_data_df.rename(columns={'yr_renovated': 'renovated'}, inplace=True)
```

```
In [13]: #percentage of houses renovated
no_reno, yes_reno = (clean_house_data_df[['renovated']] > 0).value_counts()
percentage_reno = (yes_reno / (yes_reno + no_reno)) * 100
```

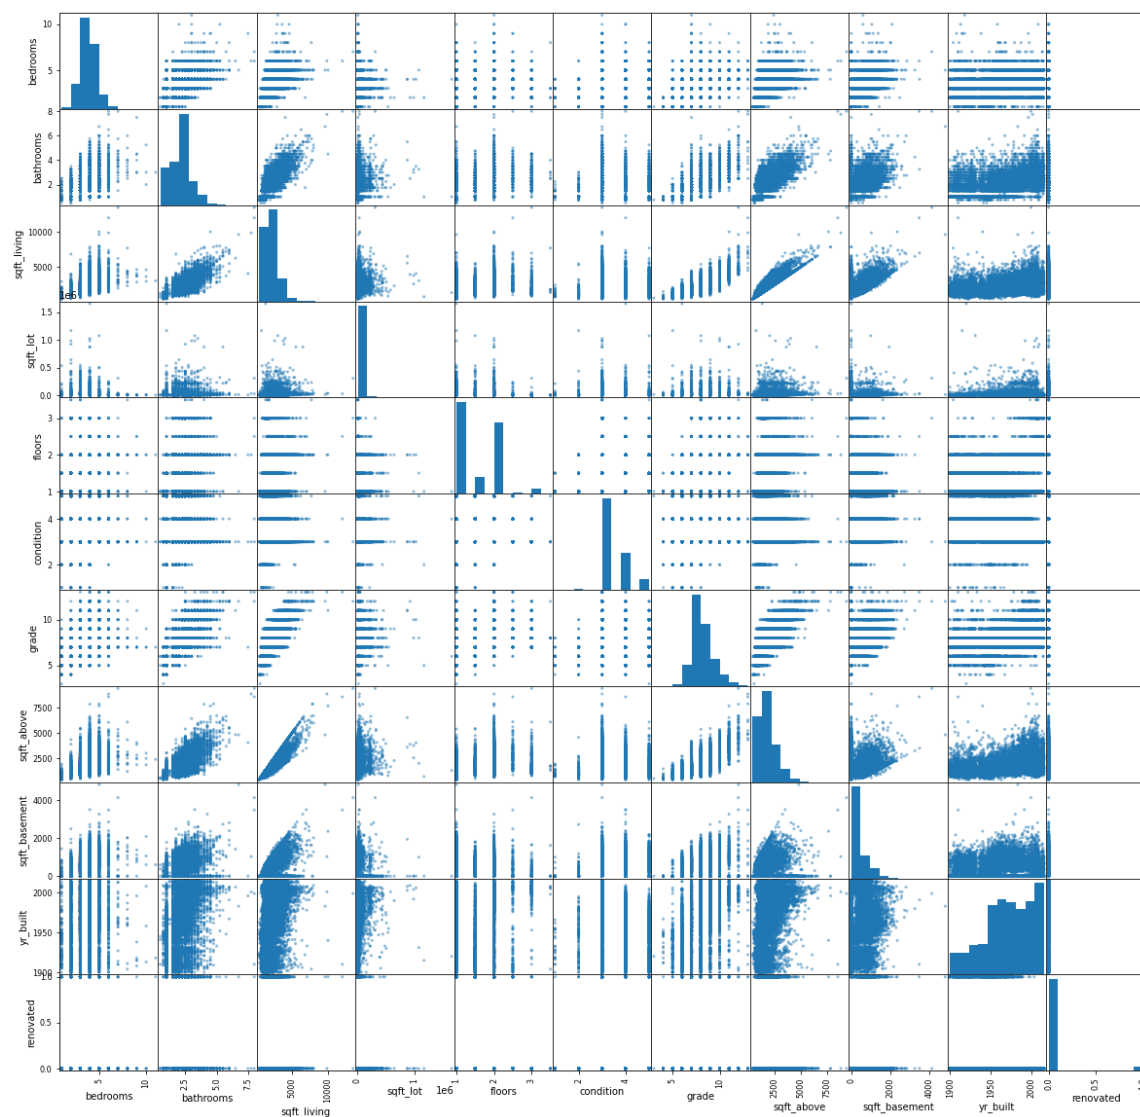
Out[13]: 4.158040027605245

Check for Multicollinearity

Variables that are highly correlated to another variable will cause problems for our regression analysis. Making the results unreliable. To fix that we look for highly correlated variables and remove some.

We removed 'sqft_basement' and 'bathrooms'


```
In [14]: house_pred = clean_house_data_df.iloc[:,2:14]
pd.plotting.scatter_matrix(house_pred, figsize = [20, 20]);
plt.show()
```



```
In [15]: house_pred.corr()
```

Out[15]:

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	condition	grade
bedrooms	1.000000	0.527457	0.592664	0.031053	0.185320	0.017068	0.366044
bathrooms	0.527457	1.000000	0.753504	0.085354	0.503906	-0.131313	0.665489
sqft_living	0.592664	0.753504	1.000000	0.167745	0.356532	-0.066163	0.764592
sqft_lot	0.031053	0.085354	0.167745	1.000000	-0.008833	-0.011021	0.108905
floors	0.185320	0.503906	0.356532	-0.008833	1.000000	-0.263477	0.460728
condition	0.017068	-0.131313	-0.066163	-0.011021	-0.263477	1.000000	-0.151378
grade	0.366044	0.665489	0.764592	0.108905	0.460728	-0.151378	1.000000
sqft_above	0.491095	0.685983	0.875885	0.175682	0.527141	-0.161946	0.758232
sqft_basement	0.310868	0.280347	0.436399	0.019582	-0.245058	0.165047	0.168573
yr_built	0.160033	0.506510	0.317275	0.048982	0.487741	-0.366227	0.445345
renovated	0.019734	0.051249	0.054029	0.003594	0.003153	-0.061755	0.016509

```
In [16]: df=house_pred.corr().abs().stack().reset_index().sort_values(0, ascending=False)

# zip the variable name columns (which were only named level_0 and level_1)
df['pairs'] = list(zip(df.level_0, df.level_1))

# set index to pairs
df.set_index(['pairs'], inplace = True)

# drop level columns
df.drop(columns=['level_1', 'level_0'], inplace = True)

# rename correlation column as cc rather than 0
df.columns = ['cc']

# drop duplicates. This could be dangerous if you have variables perfectly correlated
# for the sake of exercise, kept it in.
df.drop_duplicates(inplace=True)

df[(df.cc>0.75) & (df.cc<1)]
```

Out[16]:

	cc
pairs	
(sqft_above, sqft_living)	0.875885
(sqft_living, grade)	0.764592
(sqft_above, grade)	0.758232
(bathrooms, sqft_living)	0.753504

```
In [17]: clean_house_data_df.drop(['sqft_basement', 'bathrooms'], axis=1, inplace=True)
clean_house_data_df.reset_index(drop=True, inplace=True)
clean_house_data_df.head()
```

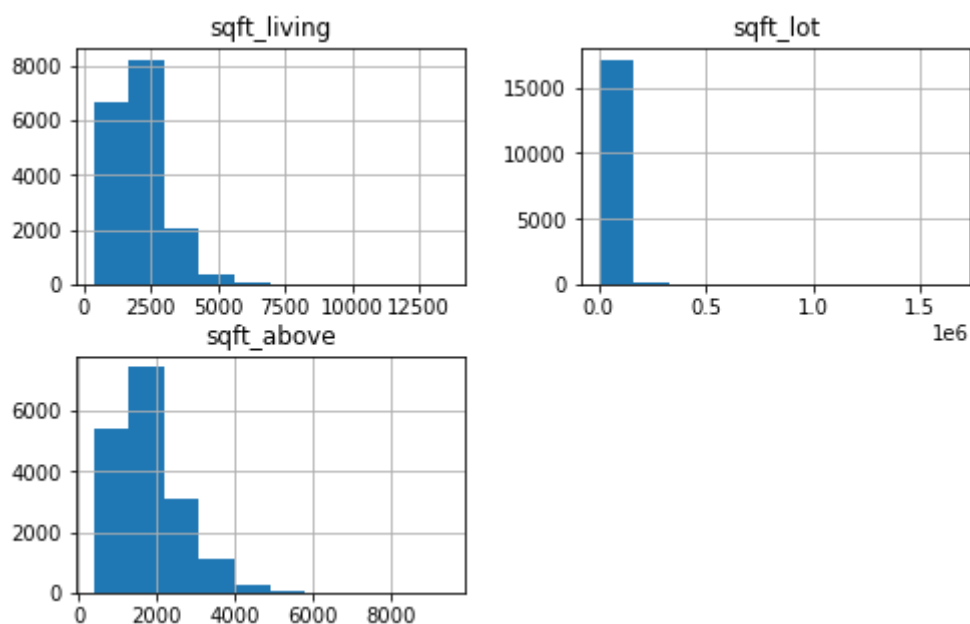
Out[17]:

	id	price	bedrooms	sqft_living	sqft_lot	floors	condition	grade	sqft_above
0	7129300520	221900.0	3	1180	5650	1.0	3	7	1180
1	6414100192	538000.0	3	2570	7242	2.0	3	7	2170
2	2487200875	604000.0	4	1960	5000	1.0	5	7	1080
3	1954400510	510000.0	3	1680	8080	1.0	3	8	1680
4	7237550310	1230000.0	4	5420	101930	1.0	3	11	3890

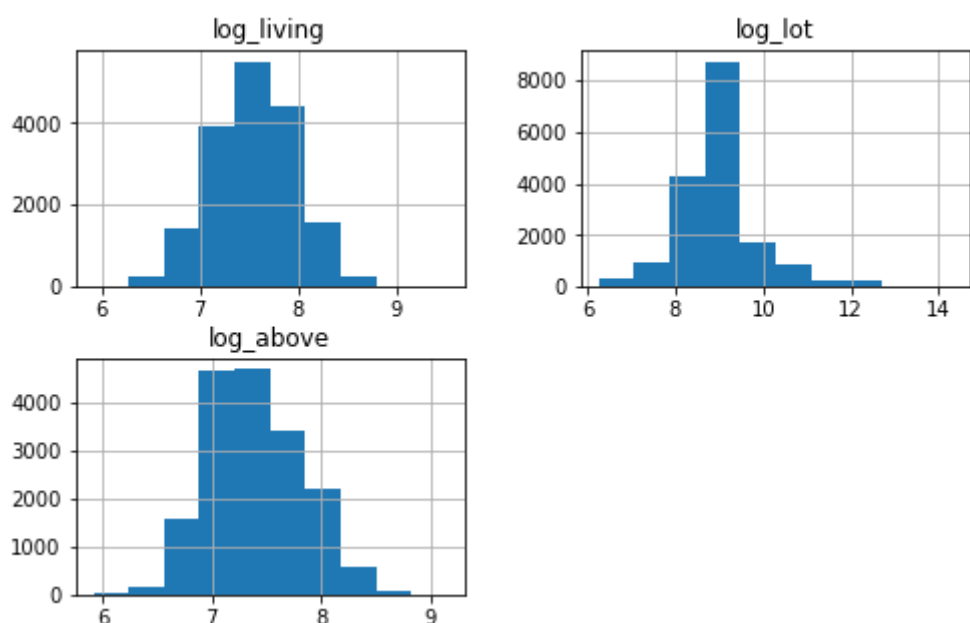
Normalise

Our data is positively skewed so we need to do log transformation to make it have a more normal distribution. After that we need to standardise our data making the mean 0

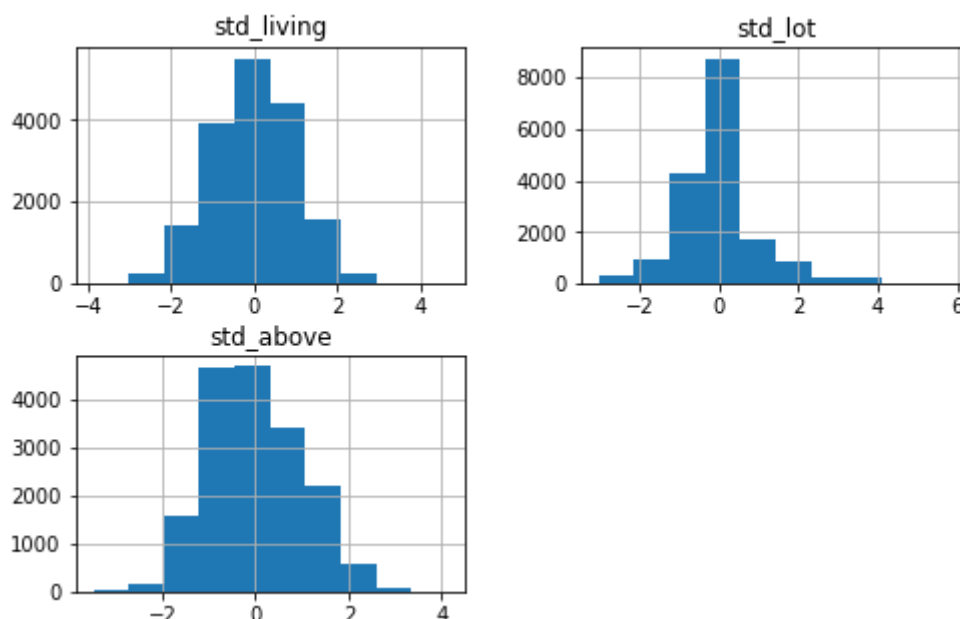
```
In [18]: # original predictors
data_pred = clean_house_data_df.iloc[:,2:10]
data_pred[['sqft_living', 'sqft_lot', 'sqft_above']].hist(figsize=[8,5]).
```



```
In [19]: # predictors after log transform
data_log = pd.DataFrame([])
data_log['log_living'] = np.log(data_pred['sqft_living'])
data_log['log_lot'] = np.log(data_pred['sqft_lot'])
data_log['log_above'] = np.log(data_pred['sqft_above'])
data_log.hist(figsize=[8,5]).
```



```
In [20]: # standard
std_scale = preprocessing.StandardScaler().fit(data_log[['log_living', 'log_
df_std = std_scale.transform(data_log[['log_living', 'log_lot', 'log_above'
standard_df = pd.DataFrame(df_std, columns=['std_living', 'std_lot', 'std_al
standard_df.hist(figsize=[8, 5]):
```



```
In [21]: clean_house_data_df2 = pd.concat([clean_house_data_df['price'], standard_df
```

One Hot Encode

For linear regression, categorical data should be transformed using one-hot encoding. In order to not have so many predictors for the year built we categorised them into 5 year increments.

```
In [22]: # categorise yr_built in 5 year increments
def categorise_yr(yr):
    yr_mod = yr % 10
    if yr_mod > 0 and yr_mod < 5:
        yr -= yr_mod
    elif yr_mod > 5 and yr_mod < 10:
        yr -= yr_mod - 5
    return yr

clean_house_data_df['yr_cat'] = clean_house_data_df['yr_built'].map(categorise_yr)
```

```
In [23]: # dummy variable
br_dum = pd.get_dummies(clean_house_data_df['bedrooms'], prefix='br', drop_
fl_dum = pd.get_dummies(clean_house_data_df['floors'], prefix='fl', drop_fi
cond_dum = pd.get_dummies(clean_house_data_df['condition'], prefix='cond',
gr_dum = pd.get_dummies(clean_house_data_df['grade'], prefix='grade', drop_
yb_dum = pd.get_dummies(clean_house_data_df['yr_cat'], prefix='yr_built', d

# add dummy variables
clean_house_data_df3 = pd.concat([clean_house_data_df2, br_dum, fl_dum, con
clean_house_data_df3.head()
```

```
Out[23]:
```

	price	std_living	std_lot	std_above	renovated	br_2	br_3	br_4	br_5	br_6	...
0	221900.0	-1.134041	-0.391525	-0.760796	0.0	0	1	0	0	0	...

	price	std_living	std_lot	std_above	renovated	br_2	br_3	br_4	br_5	br_6
1	538000.0	0.704189	-0.117255	0.667971	1.0	0	1	0	0	0
2	604000.0	0.064293	-0.526558	-1.034546	0.0	0	0	1	0	0
3	510000.0	-0.299745	0.003721	0.067739	0.0	0	1	0	0	0
4	1230000.0	2.466373	2.804408	2.036862	0.0	0	0	1	0	0

Modeling

```
In [24]: house_df = clean_house_data_df3
#variables for categorical predictors
price = house_df[['price']]
bedrooms = house_df.iloc[:,5:15]
floors = house_df.iloc[:,15:20]
condition = house_df.iloc[:,20:24]
grade = house_df.iloc[:,24:34]
yr_built = house_df.iloc[:,34:57]

house_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17388 entries, 0 to 17387
Data columns (total 57 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                  17388 non-null  float64
1   std_living             17388 non-null  float64
2   std_lot                17388 non-null  float64
3   std_above              17388 non-null  float64
4   renovated              17388 non-null  float64
5   br_2                   17388 non-null  uint8
6   br_3                   17388 non-null  uint8
7   br_4                   17388 non-null  uint8
8   br_5                   17388 non-null  uint8
9   br_6                   17388 non-null  uint8
10  br_7                   17388 non-null  uint8
11  br_8                   17388 non-null  uint8
12  br_9                   17388 non-null  uint8
13  br_10                  17388 non-null  uint8
14  br_11                  17388 non-null  uint8
15  fl_1.5                 17388 non-null  uint8
16  fl_2.0                 17388 non-null  uint8
17  fl_2.5                 17388 non-null  uint8
18  fl_3.0                 17388 non-null  uint8
19  fl_3.5                 17388 non-null  uint8
20  cond_2                 17388 non-null  uint8
21  cond_3                 17388 non-null  uint8
```

Model 1

Our first model used all the available predictors and got an R-Squared value of 0.633 which is reasonable, being able to explain 63% of variations of our model.

std_above and conditions had p-values greater than 0.05 so we will remove those for the next model.

```
In [25]: # all predictors
price = house_df['price']
predictors = house_df.iloc[:,1:57]
predictors_int = sm.add_constant(predictors)
model = sm.OLS(price, predictors_int).fit()
model.summary()
```

Out[25]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.633
Model:	OLS	Adj. R-squared:	0.632
Method:	Least Squares	F-statistic:	533.9
Date:	Sat, 11 Jun 2022	Prob (F-statistic):	0.00
Time:	21:56:41	Log-Likelihood:	-2.3890e+05
No. Observations:	17388	AIC:	4.779e+05
Df Residuals:	17331	BIC:	4.784e+05
Df Model:	56		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	6.503e+05	2.31e+05	2.815	0.005	1.98e+05	1.1e+06
std_living	1.379e+05	4276.659	32.233	0.000	1.29e+05	1.46e+05
std_lot	-1.62e+04	2316.634	-6.994	0.000	-2.07e+04	-1.17e+04
std_above	-932.2023	4489.225	-0.208	0.836	-9731.536	7867.132
renovated	8.674e+04	9159.535	9.469	0.000	6.88e+04	1.05e+05
br_2	-4.594e+04	1.93e+04	-2.375	0.018	-8.39e+04	-8032.671
br_3	-9.772e+04	1.95e+04	-5.017	0.000	-1.36e+05	-5.95e+04
br_4	-1.324e+05	2e+04	-6.617	0.000	-1.72e+05	-9.32e+04
br_5	-9.4e+04	2.11e+04	-4.458	0.000	-1.35e+05	-5.27e+04
br_6	-7.267e+04	2.53e+04	-2.871	0.004	-1.22e+05	-2.31e+04
br_7	5526.2010	4.96e+04	0.111	0.911	-9.17e+04	1.03e+05
br_8	1.745e+05	7.15e+04	2.440	0.015	3.43e+04	3.15e+05
br_9	-1.785e+05	9.46e+04	-1.887	0.059	-3.64e+05	6923.749
br_10	-6.378e+04	1.31e+05	-0.485	0.628	-3.22e+05	1.94e+05
br_11	-3.947e+05	2.26e+05	-1.746	0.081	-8.38e+05	4.84e+04
fl_1.5	-1.217e+04	7435.389	-1.636	0.102	-2.67e+04	2408.040
fl_2.0	7148.2418	6289.004	1.137	0.256	-5178.840	1.95e+04
fl_2.5	1.036e+05	2.11e+04	4.909	0.000	6.22e+04	1.45e+05
fl_3.0	8.658e+04	1.27e+04	6.817	0.000	6.17e+04	1.11e+05
fl_3.5	1.178e+05	9.29e+04	1.269	0.205	-6.42e+04	3e+05
cond_2	-5.008e+04	5.17e+04	-0.969	0.333	-1.51e+05	5.13e+04
cond_3	-3.032e+04	4.84e+04	-0.627	0.531	-1.25e+05	6.45e+04
cond_4	1384.9881	4.84e+04	0.029	0.977	-9.35e+04	9.63e+04
cond_5	4.229e+04	4.87e+04	0.869	0.385	-5.31e+04	1.38e+05
grade_4	-3335.2368	2.31e+05	-0.014	0.988	-4.57e+05	4.5e+05
grade_5	-5.714e+04	2.26e+05	-0.253	0.800	-5e+05	3.86e+05
grade_6	-3.668e+04	2.26e+05	-0.162	0.871	-4.79e+05	4.06e+05
grade_7	4.634e+04	2.26e+05	0.205	0.837	-3.96e+05	4.89e+05
grade_8	1.545e+05	2.26e+05	0.684	0.494	-2.88e+05	5.97e+05
grade_9	3.306e+05	2.26e+05	1.463	0.144	-1.12e+05	7.74e+05
grade_10	5.889e+05	2.26e+05	2.604	0.009	1.46e+05	1.03e+06
grade_11	9.444e+05	2.26e+05	4.171	0.000	5.01e+05	1.39e+06
grade_12	1.632e+06	2.28e+05	7.165	0.000	1.19e+06	2.08e+06
grade_13	2.823e+06	2.36e+05	11.949	0.000	2.36e+06	3.29e+06
yr_built_1905	2.457e+04	2.07e+04	1.188	0.235	-1.6e+04	6.51e+04
yr_built_1910	7147.3649	2.08e+04	0.344	0.731	-3.36e+04	4.79e+04
yr_built_1915	-799.0763	2.08e+04	-0.038	0.969	-4.16e+04	4e+04
yr_built_1920	-1441.6611	2.02e+04	-0.071	0.943	-4.1e+04	3.81e+04

yr_built_1925	8021.9628	1.92e+04	0.418	0.676	-2.96e+04	4.56e+04
yr_built_1930	-3.122e+04	2.34e+04	-1.334	0.182	-7.71e+04	1.46e+04
yr_built_1935	6773.2950	2.24e+04	0.302	0.763	-3.72e+04	5.07e+04
yr_built_1940	-4.195e+04	1.9e+04	-2.206	0.027	-7.92e+04	-4675.008
yr_built_1945	-4.731e+04	1.88e+04	-2.515	0.012	-8.42e+04	-1.04e+04
yr_built_1950	-5.598e+04	1.84e+04	-3.041	0.002	-9.21e+04	-1.99e+04
yr_built_1955	-1.168e+05	1.85e+04	-6.296	0.000	-1.53e+05	-8.04e+04
yr_built_1960	-1.443e+05	1.86e+04	-7.740	0.000	-1.81e+05	-1.08e+05
yr_built_1965	-1.639e+05	1.84e+04	-8.892	0.000	-2e+05	-1.28e+05
yr_built_1970	-1.5e+05	1.96e+04	-7.636	0.000	-1.88e+05	-1.11e+05
yr_built_1975	-2.1e+05	1.84e+04	-11.437	0.000	-2.46e+05	-1.74e+05
yr_built_1980	-1.811e+05	1.9e+04	-9.513	0.000	-2.18e+05	-1.44e+05
yr_built_1985	-2.512e+05	1.87e+04	-13.431	0.000	-2.88e+05	-2.15e+05
yr_built_1990	-2.949e+05	1.89e+04	-15.567	0.000	-3.32e+05	-2.58e+05
yr_built_1995	-2.69e+05	1.92e+04	-14.040	0.000	-3.07e+05	-2.31e+05
yr_built_2000	-2.437e+05	1.87e+04	-13.058	0.000	-2.8e+05	-2.07e+05
yr_built_2005	-2.308e+05	1.85e+04	-12.482	0.000	-2.67e+05	-1.95e+05
yr_built_2010	-2.254e+05	1.91e+04	-11.810	0.000	-2.63e+05	-1.88e+05
yr_built_2015	-1.922e+05	4.61e+04	-4.167	0.000	-2.83e+05	-1.02e+05

Omnibus:	13866.889	Durbin-Watson:	1.976
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1113381.178
Skew:	3.275	Prob(JB):	0.00
Kurtosis:	41.650	Cond. No.	685.

Notes:

Model 2

Removing the two predictors have lowered our R-Squared score slightly

Skew and Kurtosis is still quite high

Doing a QQ-plot shows it is not normal so we run log transformation on price as well.

```
In [26]: data1 = house_df[['std_living', 'renovated']]
data2 = pd.concat([bedrooms, floors, grade, yr_built], axis=1)
predictors = pd.concat([data1, data2], axis=1)

predictors_int = sm.add_constant(predictors)
model = sm.OLS(price, predictors_int).fit()
model.summary()
```

Out[26]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.629
-----------------------	-------	-------------------	-------

Model: OLS **Adj. R-squared:** 0.628
Method: Least Squares **F-statistic:** 587.8
Date: Sat, 11 Jun 2022 **Prob (F-statistic):** 0.00
Time: 21:56:42 **Log-Likelihood:** -2.3900e+05
No. Observations: 17388 **AIC:** 4.781e+05
Df Residuals: 17337 **BIC:** 4.785e+05
Df Model: 50
Covariance Type: nonrobust

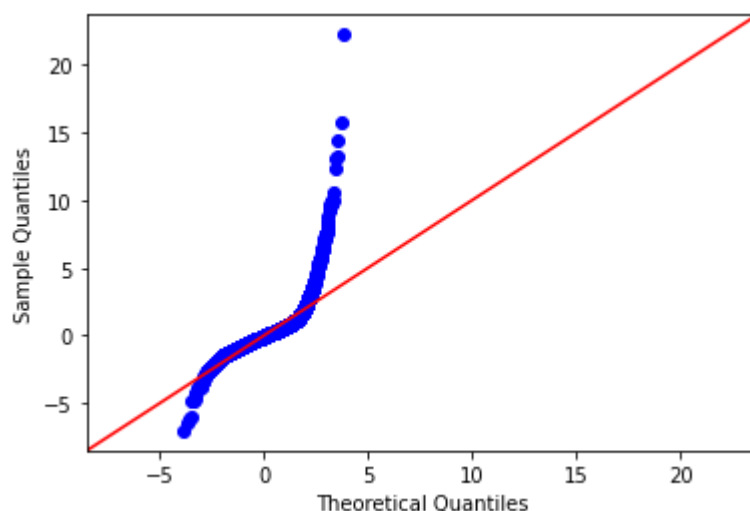
	coef	std err	t	P> t	[0.025	0.975]
const	6.745e+05	2.27e+05	2.972	0.003	2.3e+05	1.12e+06
std_living	1.328e+05	3277.778	40.514	0.000	1.26e+05	1.39e+05
renovated	6.815e+04	9054.610	7.527	0.000	5.04e+04	8.59e+04
br_2	-3.806e+04	1.94e+04	-1.960	0.050	-7.61e+04	0.450
br_3	-8.783e+04	1.95e+04	-4.497	0.000	-1.26e+05	-4.95e+04
br_4	-1.212e+05	2.01e+04	-6.041	0.000	-1.61e+05	-8.19e+04
br_5	-8.199e+04	2.11e+04	-3.877	0.000	-1.23e+05	-4.05e+04
br_6	-6.271e+04	2.54e+04	-2.470	0.014	-1.12e+05	-1.29e+04
br_7	1.683e+04	4.98e+04	0.338	0.735	-8.08e+04	1.14e+05
br_8	1.908e+05	7.19e+04	2.655	0.008	4.99e+04	3.32e+05
br_9	-1.843e+05	9.5e+04	-1.939	0.052	-3.71e+05	1971.127
br_10	-4.838e+04	1.32e+05	-0.366	0.714	-3.07e+05	2.11e+05
br_11	-3.934e+05	2.27e+05	-1.731	0.083	-8.39e+05	5.2e+04
fl_1.5	-1.181e+04	7199.974	-1.641	0.101	-2.59e+04	2300.042
fl_2.0	1.233e+04	5608.036	2.198	0.028	1335.255	2.33e+04
fl_2.5	1.146e+05	2.09e+04	5.480	0.000	7.36e+04	1.56e+05
fl_3.0	1.09e+05	1.21e+04	8.980	0.000	8.52e+04	1.33e+05
fl_3.5	1.475e+05	9.33e+04	1.581	0.114	-3.53e+04	3.3e+05
grade_4	-5.935e+04	2.32e+05	-0.255	0.798	-5.15e+05	3.96e+05
grade_5	-9.82e+04	2.27e+05	-0.432	0.666	-5.44e+05	3.47e+05
grade_6	-7.345e+04	2.27e+05	-0.324	0.746	-5.18e+05	3.71e+05
grade_7	1.701e+04	2.27e+05	0.075	0.940	-4.28e+05	4.62e+05
grade_8	1.262e+05	2.27e+05	0.556	0.578	-3.19e+05	5.71e+05
grade_9	2.978e+05	2.27e+05	1.311	0.190	-1.47e+05	7.43e+05
grade_10	5.519e+05	2.27e+05	2.429	0.015	1.07e+05	9.97e+05
grade_11	9.04e+05	2.28e+05	3.973	0.000	4.58e+05	1.35e+06
grade_12	1.587e+06	2.29e+05	6.935	0.000	1.14e+06	2.04e+06
grade_13	2.781e+06	2.37e+05	11.717	0.000	2.32e+06	3.25e+06
yr_built_1905	2.773e+04	2.08e+04	1.334	0.182	-1.3e+04	6.85e+04
yr_built_1910	8816.9599	2.09e+04	0.422	0.673	-3.22e+04	4.98e+04

yr_built_1915	-626.7098	2.09e+04	-0.030	0.976	-4.16e+04	4.03e+04
yr_built_1920	-4932.6092	2.03e+04	-0.243	0.808	-4.47e+04	3.48e+04
yr_built_1925	7723.2561	1.93e+04	0.401	0.689	-3e+04	4.55e+04
yr_built_1930	-3.892e+04	2.35e+04	-1.657	0.098	-8.5e+04	7122.242
yr_built_1935	-5119.2584	2.25e+04	-0.227	0.820	-4.92e+04	3.9e+04
yr_built_1940	-5.215e+04	1.91e+04	-2.734	0.006	-8.95e+04	-1.48e+04
yr_built_1945	-6.156e+04	1.89e+04	-3.265	0.001	-9.85e+04	-2.46e+04
yr_built_1950	-7.276e+04	1.84e+04	-3.947	0.000	-1.09e+05	-3.66e+04
yr_built_1955	-1.348e+05	1.86e+04	-7.269	0.000	-1.71e+05	-9.85e+04
yr_built_1960	-1.622e+05	1.86e+04	-8.705	0.000	-1.99e+05	-1.26e+05
yr_built_1965	-1.808e+05	1.84e+04	-9.818	0.000	-2.17e+05	-1.45e+05
yr_built_1970	-1.699e+05	1.96e+04	-8.660	0.000	-2.08e+05	-1.31e+05
yr_built_1975	-2.334e+05	1.83e+04	-12.741	0.000	-2.69e+05	-1.97e+05
yr_built_1980	-2.125e+05	1.89e+04	-11.218	0.000	-2.5e+05	-1.75e+05
yr_built_1985	-2.85e+05	1.86e+04	-15.322	0.000	-3.21e+05	-2.49e+05
yr_built_1990	-3.329e+05	1.88e+04	-17.694	0.000	-3.7e+05	-2.96e+05
yr_built_1995	-3.075e+05	1.9e+04	-16.151	0.000	-3.45e+05	-2.7e+05
yr_built_2000	-2.769e+05	1.86e+04	-14.925	0.000	-3.13e+05	-2.41e+05
yr_built_2005	-2.579e+05	1.84e+04	-14.006	0.000	-2.94e+05	-2.22e+05
yr_built_2010	-2.537e+05	1.9e+04	-13.361	0.000	-2.91e+05	-2.16e+05
yr_built_2015	-2.171e+05	4.63e+04	-4.688	0.000	-3.08e+05	-1.26e+05

Omnibus:	13776.736	Durbin-Watson:	1.976
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1075858.613
Skew:	3.250	Prob(JB):	0.00
Kurtosis:	40.983	Cond. No.	600.

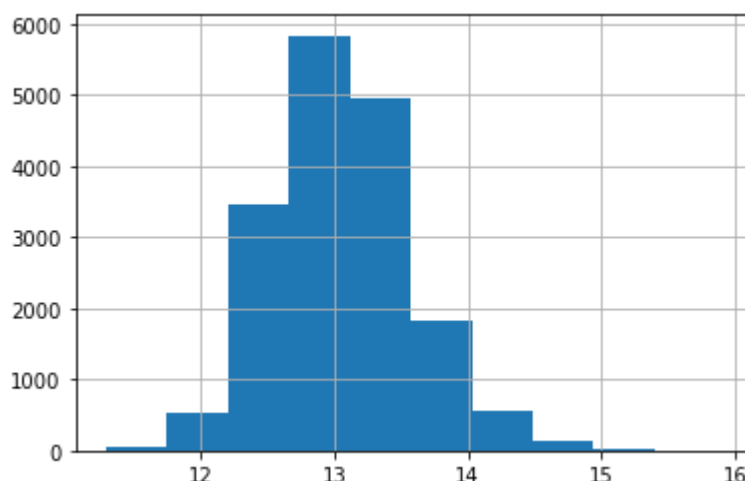
Notes:

```
In [27]: resid1 = model.resid
fig = sm.graphics.gcf().gcf().add_subplot(1, 1, 1, dist=stats.norm, line='45', fit=True)
```



```
In [28]: price = np.log(house_df['price'])
price.hist()
```

Out[28]: <AxesSubplot:>



Model 3

After log transformation on price we can see if has improved the distribution to be more normal

Improved Skewness from highly positive skew to slightly negative skew. Improved skew which is now between -0.5 and 0.5 meaning the data is pretty symmetrical as shown in the QQ-plot below

R-Squared value has also increased to 0.646 meaning 64.6% of the variance is explained by the model.

```
In [29]: data1 = house_df[['std_living', 'renovated']]
data2 = pd.concat([bedrooms, floors, grade, yr_built], axis=1)
predictors = pd.concat([data1, data2], axis=1)

predictors_int = sm.add_constant(predictors)
model = sm.OLS(price, predictors_int).fit()
model.summary()
```

Out[29]:

OLS Regression Results

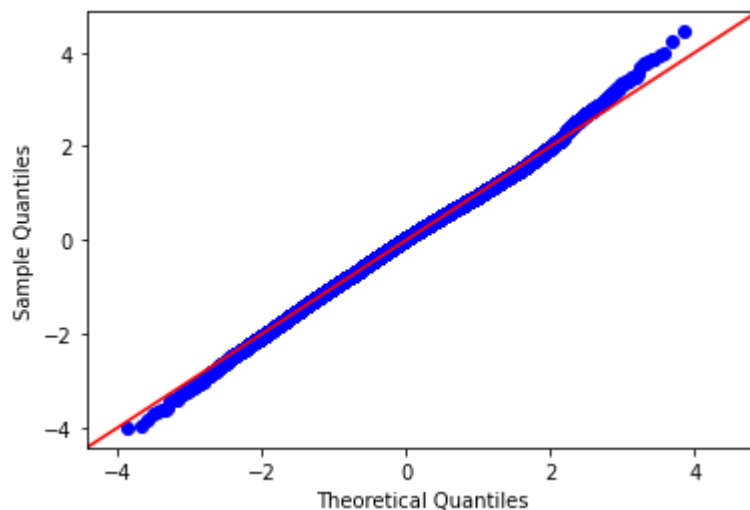
Dep. Variable:	price	R-squared:	0.646
Model:	OLS	Adj. R-squared:	0.645
Method:	Least Squares	F-statistic:	633.8
Date:	Sat, 11 Jun 2022	Prob (F-statistic):	0.00
Time:	21:56:42	Log-Likelihood:	-4467.5
No. Observations:	17388	AIC:	9037.
Df Residuals:	17337	BIC:	9433.
Df Model:	50		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	13.2076	0.315	41.941	0.000	12.590	13.825
std_living	0.2268	0.005	49.874	0.000	0.218	0.236
renovated	0.0344	0.013	2.735	0.006	0.010	0.059
br_2	-0.0489	0.027	-1.817	0.069	-0.102	0.004
br_3	-0.1369	0.027	-5.050	0.000	-0.190	-0.084
br_4	-0.1696	0.028	-6.094	0.000	-0.224	-0.115
br_5	-0.1508	0.029	-5.139	0.000	-0.208	-0.093
br_6	-0.1633	0.035	-4.636	0.000	-0.232	-0.094
br_7	-0.1575	0.069	-2.279	0.023	-0.293	-0.022
br_8	0.0288	0.100	0.289	0.773	-0.167	0.224
br_9	-0.2198	0.132	-1.667	0.096	-0.478	0.039
br_10	-0.0326	0.183	-0.178	0.859	-0.392	0.327
br_11	-0.4777	0.315	-1.515	0.130	-1.096	0.140
fl_1.5	0.0006	0.010	0.056	0.955	-0.019	0.020
fl_2.0	0.0139	0.008	1.788	0.074	-0.001	0.029
fl_2.5	0.0303	0.029	1.043	0.297	-0.027	0.087
fl_3.0	0.1887	0.017	11.199	0.000	0.156	0.222
fl_3.5	0.1056	0.129	0.816	0.415	-0.148	0.359
grade_4	-0.3337	0.322	-1.035	0.301	-0.966	0.298
grade_5	-0.2638	0.315	-0.837	0.403	-0.882	0.354
grade_6	-0.0925	0.315	-0.294	0.769	-0.710	0.525
grade_7	0.1851	0.315	0.588	0.557	-0.432	0.802
grade_8	0.4256	0.315	1.351	0.177	-0.192	1.043
grade_9	0.6683	0.315	2.121	0.034	0.051	1.286
grade_10	0.9005	0.315	2.856	0.004	0.283	1.519
grade_11	1.1192	0.316	3.545	0.000	0.500	1.738
grade_12	1.3894	0.318	4.376	0.000	0.767	2.012

grade_13	1.6329	0.329	4.958	0.000	0.987	2.278
yr_built_1905	0.0069	0.029	0.239	0.811	-0.050	0.063
yr_built_1910	-0.0545	0.029	-1.879	0.060	-0.111	0.002
yr_built_1915	-0.0442	0.029	-1.526	0.127	-0.101	0.013
yr_built_1920	-0.0353	0.028	-1.253	0.210	-0.090	0.020
yr_built_1925	-0.0489	0.027	-1.831	0.067	-0.101	0.003
yr_built_1930	-0.1559	0.033	-4.781	0.000	-0.220	-0.092
yr_built_1935	-0.1078	0.031	-3.451	0.001	-0.169	-0.047
yr_built_1940	-0.1731	0.026	-6.542	0.000	-0.225	-0.121
yr_built_1945	-0.1642	0.026	-6.277	0.000	-0.215	-0.113
yr_built_1950	-0.2182	0.026	-8.529	0.000	-0.268	-0.168
yr_built_1955	-0.3376	0.026	-13.114	0.000	-0.388	-0.287
yr_built_1960	-0.3906	0.026	-15.107	0.000	-0.441	-0.340
yr_built_1965	-0.4303	0.026	-16.840	0.000	-0.480	-0.380
yr_built_1970	-0.3855	0.027	-14.160	0.000	-0.439	-0.332
yr_built_1975	-0.4821	0.025	-18.968	0.000	-0.532	-0.432
yr_built_1980	-0.4429	0.026	-16.848	0.000	-0.494	-0.391
yr_built_1985	-0.5286	0.026	-20.481	0.000	-0.579	-0.478
yr_built_1990	-0.6020	0.026	-23.060	0.000	-0.653	-0.551
yr_built_1995	-0.5464	0.026	-20.687	0.000	-0.598	-0.495
yr_built_2000	-0.5149	0.026	-20.004	0.000	-0.565	-0.464
yr_built_2005	-0.4995	0.026	-19.551	0.000	-0.550	-0.449
yr_built_2010	-0.5064	0.026	-19.218	0.000	-0.558	-0.455
yr_built_2015	-0.4373	0.064	-6.806	0.000	-0.563	-0.311

Omnibus:	62.741	Durbin-Watson:	1.968
Prob(Omnibus):	0.000	Jarque-Bera (JB):	82.237
Skew:	-0.039	Prob(JB):	1.39e-18
Kurtosis:	3.328	Cond. No.	600.

```
In [30]: resid1 = model.resid
fig = sm.graphics.gcf().suptitle('Residuals vs Predictions', line='45', fit=True)
```

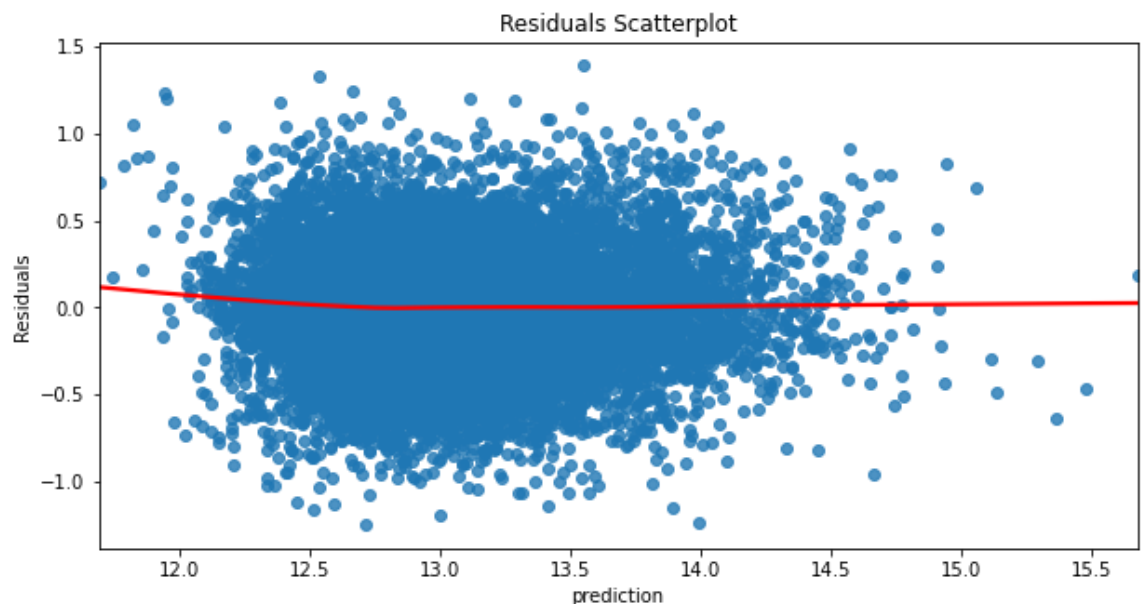


Homoscedasticity

Scatterplot to show homoscedasticity. No cone like pattern

```
In [31]: plt.figure(figsize=(10,5))
sns.regplot(x=model.predict(), y=model.resid, lowess=True, line_kws={'color': 'red'})
plt.title('Residuals Scatterplot')
plt.xlabel('prediction')
plt.ylabel('Residuals')
```

```
Out[31]: Text(0, 0.5, 'Residuals')
```



Training

With training and test MSE being similar, we can expect the model to perform similarly on different data.

Accuracy of the model is 63.42

```
In [32]: y = price
```

```

X = predictors
X_train, X_test, y_train, y_test = train_test_split(X, y)

linreg = LinearRegression()
linreg.fit(X_train, y_train)

y_hat_train = linreg.predict(X_train)
y_hat_test = linreg.predict(X_test)

train_mse = mean_squared_error(y_train, y_hat_train)
test_mse = mean_squared_error(y_test, y_hat_test)
print('Train Mean Squared Error:', train_mse)
print('Test Mean Squared Error:', test_mse)

```

Train Mean Squared Error: 0.09779224976352793
 Test Mean Squared Error: 0.09848355578437429

```

In [33]: from sklearn.metrics import r2_score

Accuracy=r2_score(y_test,y_hat_test)*100
print("Accuracy of the model is % 2f" %Accuracy)

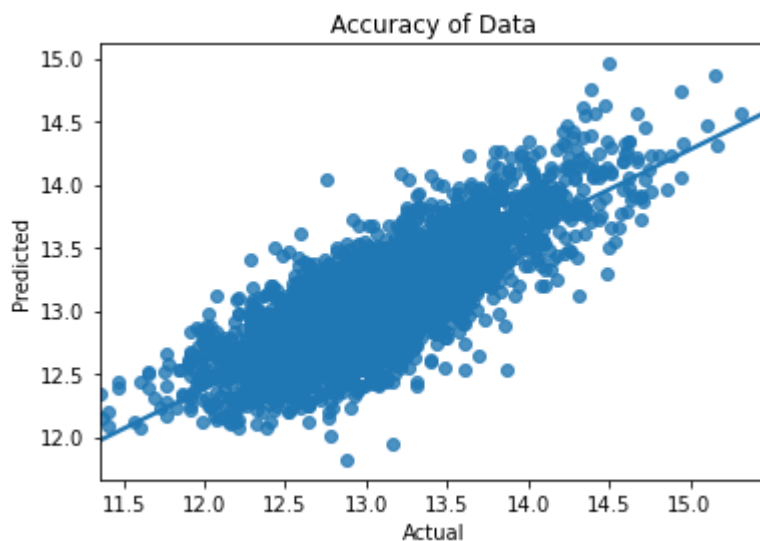
Accuracy of the model is 63.10

```

```

In [34]: plot = sns.regplot(x=y_test, y=y_hat_test, ci=None) set(title='Accuracy of Data'

```



Conclusion

With our final model the OLS regression results tell us that the R-Squared value is 0.646 meaning 64.6% of the variance can be explained by the model. The results also tells us the skewness is -0.039 which is between -0.5 to 0.5 meaning the data is symmetrical, satisfying the normality assumption. This can also be seen from the QQ-plot with points mostly following the line. If we observed a QQ-plot like in model 2 then the distribution would be non-normal. Another assumption for linear regression is that data must be homoscedastic. To check this we created a scatterplot and did not observe any cone like shapes which would indicated the data is heteroscedastic.

The living space of a property has the strongest relationship with house prices. This is determined by the t value of 49.874 which tells us how statistically significant the coefficient is. This makes sense as we spend most of the time inside the house and having a larger living area generally means more rooms or floors making it appealing to buyers. Floor had

the next highest t value of 11.1 but only for 3 floors. Having more floors generally means more living area which can increase the value. Grades are also significant as it represents the quality of the home. Renovated has a t value of 2.7 which is not as high. This might be because the data for unrenovated homes heavily outweighed that of renovated homes.